



## Basic Imports

```
1 import tensorflow as tf
2 import numpy as np
3 import os
4 import cv2
5 import matplotlib.pyplot as plt
6 from tensorflow.examples.tutorials.mnist import input_data
```

### Importing the MNIST dataset

```
1 mnist = tf.keras.datasets.mnist
2
3 (x_train, y_train), (x_test, y_test) = mnist.load_data()
4
5 print(type(x_train))
6 x_train = x_train.astype('float32')
7 x_train /= 255
8 x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
9 x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)
10 # y_train = y_train.reshape(y_train.shape[0],1)
11
12 # splitSz = x_train.shape[0]
13 # splitSz = int(1.0*splitSz)
14
15 # print(y_train.shape)
16 # x_train, x_test = x_train[:splitSz,:], x_train[splitSz:,:]
17 # y_train, y_test = y_train[:splitSz], y_train[splitSz:]
18
19 print(y_train.shape)
20 print(y_train[:10])
21 print(x_test.shape)
```

```
↳ <class 'numpy.ndarray'>
(60000,)
[5 0 4 1 9 2 1 3 1 4]
(10000, 28, 28, 1)
```

## ▼ Model

1. 7x7 Convolutional Layer with 32 filters and stride of 1.
2. ReLU Activation Layer.
3. Batch Normalization Layer
4. 2x2 Max Pooling layer with a stride of 2
5. fully connected layer with 1024 output units.
6. ReLU Activation Layer.

```
1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(32, (3, 3), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.BatchNormalization())
4 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
5 model.add(tf.keras.layers.Flatten())
6 model.add(tf.keras.layers.Dense(1024, activation="relu"))
7
```

```
8 model.add(tf.keras.layers.Dense(10, activation='softmax'))  
9  
10 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])  
11 model.summary()
```

↳ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/r  
Instructions for updating:  
Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	320
=====		
batch_normalization_v1 (BatchNormalization)	(None, 28, 28, 32)	128
=====		
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
=====		
flatten (Flatten)	(None, 6272)	0
=====		
dense (Dense)	(None, 1024)	6423552
=====		
dense_1 (Dense)	(None, 10)	10250
=====		
Total params: 6,434,250		
Trainable params: 6,434,186		
Non-trainable params: 64		

```
1 print(y_train.shape)  
2  
3 model.fit(x_train, y_train, batch_size=64, epochs=10)
```

↳ (60000,)  
Epoch 1/10  
60000/60000 [=====] - 226s 4ms/sample - loss: 0.2494 - acc: 0  
Epoch 2/10  
60000/60000 [=====] - 226s 4ms/sample - loss: 0.0396 - acc: 0  
Epoch 3/10  
60000/60000 [=====] - 226s 4ms/sample - loss: 0.0235 - acc: 0  
Epoch 4/10  
60000/60000 [=====] - 226s 4ms/sample - loss: 0.0178 - acc: 0  
Epoch 5/10  
60000/60000 [=====] - 227s 4ms/sample - loss: 0.0176 - acc: 0  
Epoch 6/10  
60000/60000 [=====] - 226s 4ms/sample - loss: 0.0135 - acc: 0  
Epoch 7/10  
60000/60000 [=====] - 227s 4ms/sample - loss: 0.0104 - acc: 0  
Epoch 8/10  
60000/60000 [=====] - 230s 4ms/sample - loss: 0.0123 - acc: 0  
Epoch 9/10  
60000/60000 [=====] - 227s 4ms/sample - loss: 0.0089 - acc: 0  
Epoch 10/10  
60000/60000 [=====] - 227s 4ms/sample - loss: 0.0089 - acc: 0  
<tensorflow.python.keras.callbacks.History at 0x7f0b32470748>

```
1 score = model.evaluate(x_test, y_test)
2 print('\n', 'Test Acc = ', score[1])
```

```
⇒ 10000/10000 [=====] - 8s 805us/sample - loss: 4.7237 - acc: 0
```

```
Test Acc = 0.7065
```

```
1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(32, (3, 3), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.BatchNormalization())
4 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
5 model.add(tf.keras.layers.Flatten())
6 model.add(tf.keras.layers.Dense(1024, activation="relu"))
7 model.add(tf.keras.layers.Dense(1024, activation="relu"))
8
9 model.add(tf.keras.layers.Dense(10, activation='softmax'))
10
11 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
12 model.summary()
13
14 print(y_train.shape)
15
16 model.fit(x_train, y_train, batch_size=64, epochs=5)
17
18 score = model.evaluate(x_test, y_test)
19 print('\n', 'Test Acc = ', score[1])
```

```
⇒
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.BatchNormalization())
4 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
5 model.add(tf.keras.layers.Flatten())
6 model.add(tf.keras.layers.Dense(1024, activation="relu"))
7
8 model.add(tf.keras.layers.Dense(10, activation='softmax'))
9
10 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
11 model.summary()
12
13 print(y_train.shape)
14
15 model.fit(x_train, y_train, batch_size=64, epochs=5)
16
17 score = model.evaluate(x_test, y_test)
18 print('\n', 'Test Acc = ', score[1])

```

→

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 28, 28, 32)	1600
batch_normalization_v1_2 (BatchNormalization)	(None, 28, 28, 32)	128
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
flatten_2 (Flatten)	(None, 6272)	0
dense_5 (Dense)	(None, 1024)	6423552
dense_6 (Dense)	(None, 10)	10250

Total params: 6,435,530  
Trainable params: 6,435,466  
Non-trainable params: 64

---

(60000,)  
Epoch 1/5  
60000/60000 [=====] - 267s 4ms/sample - loss: 0.2479 - acc: 0  
Epoch 2/5  
60000/60000 [=====] - 268s 4ms/sample - loss: 0.0415 - acc: 0  
Epoch 3/5  
60000/60000 [=====] - 268s 4ms/sample - loss: 0.0266 - acc: 0  
Epoch 4/5  
60000/60000 [=====] - 267s 4ms/sample - loss: 0.0235 - acc: 0  
Epoch 5/5  
60000/60000 [=====] - 266s 4ms/sample - loss: 0.0196 - acc: 0  
10000/10000 [=====] - 12s 1ms/sample - loss: 0.3357 - acc: 0.

Test Acc = 0.9791

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(64, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.BatchNormalization())
4 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
5 model.add(tf.keras.layers.Flatten())
6 model.add(tf.keras.layers.Dense(1024, activation="relu"))
7
8 model.add(tf.keras.layers.Dense(10, activation='softmax'))
9
10 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
11 model.summary()
12
13 print(y_train.shape)
14
15 model.fit(x_train, y_train, batch_size=64, epochs=5)
16
17 score = model.evaluate(x_test, y_test)
18 print('\n', 'Test Acc = ', score[1])

```

→ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/r  
 Instructions for updating:  
 Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 64)	3200
batch_normalization_v1 (BatchNormalization)	(None, 28, 28, 64)	256
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1024)	12846080
dense_1 (Dense)	(None, 10)	10250
<hr/>		
Total params: 12,859,786		
Trainable params: 12,859,658		
Non-trainable params: 128		

---

(60000,)  
 Epoch 1/5  
 60000/60000 [=====] - 513s 9ms/sample - loss: 0.3795 - acc: 0  
 Epoch 2/5  
 60000/60000 [=====] - 505s 8ms/sample - loss: 0.0364 - acc: 0  
 Epoch 3/5  
 60000/60000 [=====] - 503s 8ms/sample - loss: 0.0278 - acc: 0  
 Epoch 4/5  
 60000/60000 [=====] - 507s 8ms/sample - loss: 0.0240 - acc: 0  
 Epoch 5/5  
 60000/60000 [=====] - 507s 8ms/sample - loss: 0.0205 - acc: 0  
 10000/10000 [=====] - 18s 2ms/sample - loss: 0.3191 - acc: 0.

Test Acc = 0.9801

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
4 model.add(tf.keras.layers.BatchNormalization())
5 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
6 model.add(tf.keras.layers.Flatten())
7 model.add(tf.keras.layers.Dense(1024, activation="relu"))
8
9 model.add(tf.keras.layers.Dense(10, activation='softmax'))
10
11 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
12 model.summary()
13
14 print(y_train.shape)
15
16 model.fit(x_train, y_train, batch_size=64, epochs=5)
17
18 score = model.evaluate(x_test, y_test)
19 print('\n', 'Test Acc = ', score[1])

```



Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 28, 28, 32)	1600
conv2d_2 (Conv2D)	(None, 28, 28, 32)	50208
batch_normalization_v1_1 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_2 (Dense)	(None, 1024)	6423552
dense_3 (Dense)	(None, 10)	10250
<hr/>		
Total params: 6,485,738		
Trainable params: 6,485,674		
Non-trainable params: 64		

---

(60000,)

Epoch 1/5  
60000/60000 [=====] - 806s 13ms/sample - loss: 0.1211 - acc: 0.9778

Epoch 2/5  
60000/60000 [=====] - 783s 13ms/sample - loss: 0.0459 - acc: 0.9778

Epoch 3/5  
60000/60000 [=====] - 799s 13ms/sample - loss: 0.0294 - acc: 0.9778

Epoch 4/5  
60000/60000 [=====] - 791s 13ms/sample - loss: 0.0268 - acc: 0.9778

Epoch 5/5  
60000/60000 [=====] - 787s 13ms/sample - loss: 0.0212 - acc: 0.9778

10000/10000 [=====] - 33s 3ms/sample - loss: 0.3530 - acc: 0.9778

Test Acc = 0.9778

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.BatchNormalization())

```

```

4 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
5 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
6 model.add(tf.keras.layers.Flatten())
7 model.add(tf.keras.layers.Dense(1024, activation="relu"))
8
9 model.add(tf.keras.layers.Dense(10, activation='softmax'))
10
11 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
12 model.summary()
13
14 print(y_train.shape)
15
16 model.fit(x_train, y_train, batch_size=64, epochs=5)
17
18 score = model.evaluate(x_test, y_test)
19 print('\n', 'Test Acc = ', score[1])

```



Layer (type)	Output Shape	Param #
<hr/>		
conv2d_3 (Conv2D)	(None, 28, 28, 32)	1600
batch_normalization_v1_2 (BatchNormalization)	(None, 28, 28, 32)	128
conv2d_4 (Conv2D)	(None, 28, 28, 32)	50208
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
flatten_2 (Flatten)	(None, 6272)	0
dense_4 (Dense)	(None, 1024)	6423552
dense_5 (Dense)	(None, 10)	10250
<hr/>		
Total params: 6,485,738		
Trainable params: 6,485,674		
Non-trainable params: 64		

---

(60000,)

Epoch 1/5

60000/60000 [=====] - 772s 13ms/sample - loss: 1.5519 - acc: 0.0000

Epoch 2/5

60000/60000 [=====] - 760s 13ms/sample - loss: 0.0421 - acc: 0.9765

Epoch 3/5

60000/60000 [=====] - 760s 13ms/sample - loss: 0.0281 - acc: 0.9765

Epoch 4/5

60000/60000 [=====] - 773s 13ms/sample - loss: 0.0223 - acc: 0.9765

Epoch 5/5

60000/60000 [=====] - 771s 13ms/sample - loss: 0.0195 - acc: 0.9765

10000/10000 [=====] - 33s 3ms/sample - loss: 0.3756 - acc: 0.9765

Test Acc = 0.9765

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(64, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.Conv2D(32, (7, 7), strides=1, activation='relu', padding="same")
4 model.add(tf.keras.layers.BatchNormalization())
5 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
6 model.add(tf.keras.layers.Flatten())

```

```

7 model.add(tf.keras.layers.Dense(1024, activation="relu"))
8 model.add(tf.keras.layers.Dense(10, activation='softmax'))
10 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
11 model.summary()
12
13 print(y_train.shape)
14
15 model.fit(x_train, y_train, batch_size=64, epochs=5)
16
17 score = model.evaluate(x_test, y_test)
18 print('\n', 'Test Acc = ', score[1])

```

→ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/r  
 Instructions for updating:  
 Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 64)	3200
conv2d_1 (Conv2D)	(None, 28, 28, 32)	100384
batch_normalization_v1 (BatchNormalization)	(None, 28, 28, 32)	128
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 1024)	6423552
dense_1 (Dense)	(None, 10)	10250

Total params: 6,537,514

Trainable params: 6,537,450

Non-trainable params: 64

---

(60000,)  
 Epoch 1/5  
 60000/60000 [=====] - 1285s 21ms/sample - loss: 0.1055 - acc: 0.0000  
 Epoch 2/5  
 60000/60000 [=====] - 1302s 22ms/sample - loss: 0.0434 - acc: 0.0000  
 Epoch 3/5  
 60000/60000 [=====] - 1298s 22ms/sample - loss: 0.0350 - acc: 0.0000  
 Epoch 4/5  
 60000/60000 [=====] - 1282s 21ms/sample - loss: 0.0287 - acc: 0.0000  
 Epoch 5/5  
 60000/60000 [=====] - 1263s 21ms/sample - loss: 0.0213 - acc: 0.0000  
 10000/10000 [=====] - 52s 5ms/sample - loss: 0.2341 - acc: 0.0000

Test Acc = 0.9853

## [Link](#)

1. input layer: [.,784]
2. layer: Conv1 -> ReLu -> MaxPool: [.,14,14,36]
3. layer: Conv2 -> ReLu -> MaxPool: [.,7,7,36]
4. layer: Conv3 -> ReLu -> MaxPool: [.,4,4,36]
5. layer: FC -> ReLu: [.,576]
6. output layer: FC -> ReLu: [.,10]

```
1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Conv2D(36, (7, 7), strides=1, activation='relu', padding="same")
3 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
4 model.add(tf.keras.layers.Conv2D(36, (7, 7), strides=1, activation='relu', padding="same")
5 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
6 model.add(tf.keras.layers.Conv2D(36, (7, 7), strides=1, activation='relu', padding="same")
7 model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
8 # model.add(tf.keras.layers.BatchNormalization())
9 # model.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
10 model.add(tf.keras.layers.Flatten())
11 model.add(tf.keras.layers.Dense(1024, activation="relu"))
12
13 model.add(tf.keras.layers.Dense(10, activation='softmax'))
14
15 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
16 model.summary()
17
18 print(y_train.shape)
19
20 model.fit(x_train, y_train, batch_size=64, epochs=5)
21
22 score = model.evaluate(x_test, y_test)
23 print('\n', 'Test Acc = ', score[1])
```



Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 28, 28, 64)	3200
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_3 (Conv2D)	(None, 14, 14, 32)	100384
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_4 (Conv2D)	(None, 7, 7, 32)	50208
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 32)	0

```

1 model2 = tf.keras.Sequential()
2 model2.add(tf.keras.layers.Conv2D(8, (7, 7), strides=1, activation='relu', padding="same")
3 model2.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
4 # model2.add(tf.keras.layers.Conv2D(36, (7, 7), strides=1, activation='relu', padding="sa
5 # model2.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
6 # model2.add(tf.keras.layers.Conv2D(36, (7, 7), strides=1, activation='relu', padding="sa
7 # model2.add(tf.keras.layers.MaxPool2D(pool_size=(2,2), strides=2))
8 model2.add(tf.keras.layers.Flatten())
9 model2.add(tf.keras.layers.Dense(1024, activation="relu"))
10
11 model2.add(tf.keras.layers.Dense(10, activation='softmax'))
12
13 model2.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
14 model2.summary()
15
16 print(y_train.shape)
17
18 some = model2.fit(x_train, y_train, batch_size=64, epochs=1)
19
20 score = model2.evaluate(x_test, y_test)
21 pred = model2.predict(x_test)
22 # print('\n', 'Test Acc = ', score[1])
23 print(score)
24 # print('\n'," Predictions: ", pred)
25 print(some)

```

...

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 28, 28, 8)	400
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 8)	0
flatten_2 (Flatten)	(None, 1568)	0
dense_4 (Dense)	(None, 1024)	1606656
dense_5 (Dense)	(None, 10)	10250

Total params: 1,617,306

Trainable params: 1,617,306

1

(60000,)

42112/60000 [=====>.....] - ETA: 26s - loss: 0.1809